

Project Report: Intelligent Community Complaint Classification System

Author: Muhammad Affan Kabir

SAP ID: 65139

Course: Artificial Intelligence

1. Problem Statement

In modern civic management, authorities receive thousands of citizen complaints daily regarding utilities, sanitation, and economic issues. Manually sorting these unstructured text messages is labor-intensive, slow, and prone to human error. Furthermore, existing automated systems often fail to distinguish between genuine complaints and irrelevant noise (e.g., casual conversation), or they miss obvious high-priority keywords due to probabilistic uncertainty.

There is a critical need for an intelligent, automated system capable of classifying complaints with high precision while robustly filtering out irrelevant data.

2. Project Objectives

The primary objectives of this project are:

- **Automation:** To develop an AI pipeline that automatically categorizes text inputs into specific departments (Cleaning, Water, Electricity, Inflation).
- **Noise Handling:** To implement a mechanism for identifying and discarding "Irrelevant" inputs (e.g., sports, personal updates) to prevent false positives.
- **Accuracy Guarantee:** To integrate a hybrid logic system that ensures 100% classification accuracy for critical domain-specific keywords (e.g., "voltage", "sewage") regardless of context.
- **Real-time Classification:** To provide a live interface for immediate prediction and routing.

3. Dataset Description

The project utilizes a custom-generated synthetic dataset designed to simulate real-world linguistic variety while avoiding common overfitting pitfalls.

- **Source:** Generation of Synthetic dataset using a combinatorial algorithm.
- **Size: 500 Data Samples**
- **Features:**
 - **Input:** Complaint_Text (Natural language sentences).
 - **Output:** Category (Target classification).
- **Class Distribution:** The dataset is balanced across five distinct categories:
 - **Cleaning:** (e.g., "Garbage is piling up...")
 - **Water:** (e.g., "The main pipe is leaking...")
 - **Electricity:** (e.g., "The transformer sparked...")
 - **Inflation:** (e.g., "Prices of flour have doubled...")
 - **Irrelevant:** (e.g., "I am watching cricket...")
- **Context Strategy:** To prevent the model from overfitting on time-based words (e.g., assuming "yesterday" always means "Water"), a shared bank of *Common Contexts* was distributed across all categories during generation.

4. Proposed Methodology

The solution employs a **Hybrid Classification Architecture** combining deterministic rules with probabilistic Machine Learning.

A. Feature Extraction

Unstructured text is converted into numerical vectors using **TF-IDF (Term Frequency-Inverse Document Frequency)**. This technique assigns weight to words based on their importance, filtering out common stop words while highlighting domain-specific terms like "voltage" or "leak."

B. Machine Learning Model

A **Multinomial Naive Bayes (MNB)** classifier is trained on the TF-IDF vectors. MNB is chosen for its efficiency and high performance in text classification tasks where features (word counts/frequencies) follow a multinomial distribution.

How it works:

Here is the short, simple flow of how **TF-IDF** and **Naive Bayes** work together in your specific project to classify a sentence like "*The voltage is low*".

The Flow: From Text to Prediction

1. Input

You type: "**The voltage is low.**"

2. TF-IDF (*The Translator & Highlighter*)

The computer cannot read English; it only understands numbers. TF-IDF converts your sentence into a list of numbers (a vector).

- **TF (Frequency):** It counts the words.
- **IDF (Rarity):** It gives a **score** to each word based on importance.
 - "The", "is" gives **Low Score** (Common words, ignored).
 - "Voltage" gives **High Score** (Rare/Unique word, highlighted).
- **Result:** Your sentence becomes a math vector like [0.01, 0.95, 0.02, 0.5].

You are right. My previous explanation glossed over the specific mechanics of Naive Bayes too quickly.

Here is the corrected, step-by-step flow for how **Naive Bayes** actually thinks when it receives the numbers from TF-IDF.

3. Naive Bayes (*The Probability Detective*)

Once TF-IDF has highlighted the important words (like "Voltage"), Naive Bayes takes over to calculate the winner.

1. The "Naive" Setup (*Treating Words as Independent Clues*)

Naive Bayes assumes every word is a separate, independent piece of evidence. It doesn't care about grammar or word order. It just asks: "*Does this specific word appear?*"

2. Consulting the "Memory" (Likelihood)

During the **Training Phase** (Step 2 in your code), the model built a mental "probability table." It knows, for example:

- "Voltage" appears in **Electricity** complaints 95% of the time.
- "Voltage" appears in **Water** complaints 0.1% of the time.

3. The Calculation (Multiplying the Odds)

When you input "*The voltage is low*", the model calculates a score for **every** category simultaneously.

- **Checking "Electricity":**
 - Start with baseline probability (25%).
 - See word "Voltage"? \rightarrow *Multiply by huge boost (0.95)*.
 - See word "Low"? \rightarrow *Multiply by small boost (0.40)*.
 - **Result:** A very high score.
- **Checking "Water":**
 - Start with baseline probability (25%).
 - See word "Voltage"? \rightarrow *Multiply by tiny penalty (0.001)*.
 - See word "Low"? \rightarrow *Multiply by small boost (0.40)*.
 - **Result:** The score crashes to near zero because "Voltage" effectively killed the possibility.

4. The Verdict

The model compares the final scores for all 4 categories.

- **Electricity:** 0.85
- **Water:** 0.02
- **Cleaning:** 0.01
- **Inflation:** 0.01

Final Decision: The system picks **Electricity**.

Summary of the Combined Flow

1. **TF-IDF** turns text into **weighted numbers** (giving high value to "Voltage").
2. **Naive Bayes** uses those numbers to **lookup probabilities** it learned during training.

3. **Multiplication** of those probabilities determines the category with the highest likelihood.

4. Output

The system returns: **Electricity**.

C. Hybrid Prediction Logic (The "Smart" Layer)

To maximize reliability, the prediction phase follows a strict hierarchical logic:

1. **Layer 1: Rule-Based Override:** The input is first checked against a dictionary of critical keywords (e.g., *electricity*, *voltage*, *sewage*, *pipe*). If a match is found, the system bypasses the AI model and assigns the category with **100% confidence**.
2. **Layer 2: AI Classification:** If no keywords are found, the MNB model predicts the most likely category.
3. **Layer 3: Threshold Filtering:**
 - a. If the predicted category is "Irrelevant", the system outputs "Others".
 - b. If the model's confidence score is below **45% (0.45)**, the prediction is deemed unsafe, and the system outputs "Others".

5. Expected Results

The implementation is expected to deliver the following outcomes:

- **High Accuracy:** The Multinomial Naive Bayes model is expected to achieve an accuracy of **>90%** on the synthetic test set due to the clear vocabulary distinction between categories.
- **Zero-Error on Critical Terms:** Thanks to the rule-based layer, inputs containing definitive triggers like "blackout" or "leakage" or "Electricity" will never be misclassified, regardless of the surrounding sentence structure.
- **Effective Noise Filtering:** The system will correctly identify non-complaint inputs (e.g., "He is eating pizza") as "Others" rather than forcing them into a civic category.

